

# Text segmentation

## Ambiguity

Word means two different things, confusing for algorithm.

## Contextuality

Extract context from surroundings (“Are you coming? I have to study” where “I have to study” implies they are not coming).

## Multilinguality

There are a lot of different languages

## Combinatorial explosion

Rapid growth due to combinatorial consideration. For example, words can be part of multiple word classes, therefore a sentence might have two or more valid structures

## Tokenization

Creating tokens out of full text

## Word tokens

All word tokens, even duplicates

## Word types

Unique words

## Normalisation

“lowercasing”, harmonizing spelling variants (color/colour), suffix removal (wanted → want)

## Stop words

Common words to filter out, don't add much to context. “The”, “a”, etc.

## Morpheme

Smallest parts you can break up words into. One root morpheme and 0 or more affixes (draw, draw+s, draw+ing+s).

## Lexeme

All words with the same base meaning (run, ran, running).

## Lemma

The lexeme you'd find in a dictionary (run, in this case).

## Part-of-speech

Word classes. Verb, adjective, noun, proper noun, etc.

## Constituent

A part of a sentence that can be replaced with something else (“he read *the book*” → “he read *it*”).

## Syntactic head

Describes constituent (“warm water” is about water). Determines internal structure and external distribution.

## Phrase structure tree

Describes sentence structure with a tree (sentence → subj-part + verb-part). Syntax trees.

## Dependency tree

Shows what word is dependent on what by drawing arrows.

## Treebank

Describing how words fit together. Examples: phrase structure tree, dependency tree.

## Supervised machine learning

System has access to input and output.

Regression → predict numerical value given input (housing prices)

Classification → Predict which of k classes some input belongs to (parliament speeches)

## Unsupervised machine learning

Has access to input but not output.

Clustering → N/A

Topic models → Put words in classes, see how many words of each class in text. “This text is 50% sports-related”.

## Text classification

### Accuracy

Percentage of correctness.

Diagonal / whole

### Precision

“If the system predicted c, how accurate is it?”

Exact match / column

## Recall

“If the input is c, how often is the system correct?”

Exact match / row

## Naive Bayes classifier

Bag of words model

For class c:  $\text{score}(c) = P(c) * P(\text{every word in input} \mid c)$

Pick class with highest score.

## Maximum likelihood estimation

*Naive Bayes:*

- $P(c)$  = Probability of input being class c without looking at text
- $P(w \mid c)$  = Probability of word w appearing in a document of class c.

## Additive smoothing

Add k to every probability → fix issue with multiply by 0 for unknown words.

## ***Evaluate a text classifier based on accuracy, precision, and recall***

See above

## ***Apply the classification rule of the Naive Bayes classifier to a text***

Calculate score for each class, pick class with highest score.

$\text{score}(\text{class}) = P(\text{class}) * P(\text{[every word]} \mid \text{class})$

## ***Learn the probabilities of a Naive Bayes classifier using maximum likelihood estimation and additive smoothing***

$P(c) = \text{count}(\text{documents classified as } c) / \text{count}(\text{documents})$

$P(w \mid c) = \text{count}(w \text{ in documents classified as } c) / \text{count}(\text{all words in documents classified as } c).$

## Language modelling

### N-gram model

Generate sequence of words, looking N-1 words back.

next word = highest  $P(\text{[all words]} \mid n-1 \text{ previous words})$

## Maximum Likelihood Estimation

$P(w) = \text{count}(w) / \text{count}(\text{all})$

$P(w \mid u) = \text{count}(uw) / \text{count}(u) \rightarrow P(\text{“rights”} \mid \text{“your”}) = \text{count}(\text{“your rights”}) / \text{count}(\text{“your”})$

## Additive smoothing

Add  $k$  to every probability → fix issue with multiply by 0 for unknown words.

## Perplexity

$2^{\text{entropy}}$

## Entropy

Probability high or low.

Count probabilities as negative log probabilities: surprisal.

## *Learn an $n$ -gram model using additive smoothing*

$P(w) = \frac{\text{count}(w) + k}{\text{count}(\text{all}) + (k * \text{count}(\text{unique}))}$

$P(w | u) = \frac{\text{count}(uw) + k}{\text{count}(u) + (k * \text{count}(\text{unique}))}$

## *Evaluate an $n$ -gram model using entropy*

$-(1/\text{count}(\text{all})) * \log_2(P(x_1, \dots, x_N))$ .

## Part of speech tagging

### Part of speech

A category of words that play similar roles within the syntactic structure of a sentence.

### Part of speech tagging

Part of speech tagger = program that tags each word in sentence with its part of speech.

Can be approached using supervised learning (requires training data).

Ambiguity (words can have different tags) → combinatorial explosion

### Accuracy

Diagonal / whole

### Precision

Exact match / column

### Recall

Exact match / row

## Hidden Markov model (HMM)

Words have probabilities tied to each of its tags ( $\text{jag} \rightarrow \text{NN}$ ,  $\text{jag} \rightarrow \text{PN}$ )

Tags have probabilities for its next tag ( $\text{NN} \rightarrow \text{VB}$ )

HMM has two probabilities: transitional (tag2 given tag1) and output (word given tag).

Transitional first, then output at every junction.

$P(\text{VB} | \text{PN}) \rightarrow \text{amount of PN followed by VB} / \text{all occurrences of VB}$

$P(\text{jag} | \text{PN}) \rightarrow \text{amount of jag when PN} / \text{all words that are PN}$

## Multi-class perceptron

### Feature window

HMM looks back once; might want to look further, or look forward. But don't want to see too much (efficiency).

Need a feature window. Feature window sees  $x$  in front and  $x$  in back of the current word.

***Evaluate a part-of-speech tagger based on accuracy, precision, and recall***

***Compute the probability of a tagged sentence in a hidden Markov model***

Probability of tagged sentence  $\rightarrow$  product of transition and output (transition \* output)

## Syntactic Analysis (wildcard)

### Phrase structure tree

Sentence divides into parts (S  $\rightarrow$  Noun Phrase, Verb Phrase), which in turn divide into parts (NP  $\rightarrow$  Pro  $\rightarrow$  "I", VP  $\rightarrow$  Verb, NP).

### Dependency tree

"This word depends on that word". Verbs have subjects and objects, etc.

### Probabilistic context-free grammar (PCFG)

Words within sentences form phrases:

"Kim read [a book]", "Kim read [a very interesting book about grammar]"

Syntactic head  $\rightarrow$  most important word in sentence.

Context free grammar  $\rightarrow$  Phrases combine. How to combine? Context free grammar!

Example: Sentence  $\rightarrow$  NP, VB (Basically BNF)

Probabilistic  $\rightarrow$  Number of trees grows exponentially with length of sentence. Not all parse trees are relevant, only most probable.

PCFG  $\rightarrow$  Every rule  $R$  has probability  $P(R)$ , and sum of all  $P(R)$  with same left side is 1.

Tree probability = product of all  $P(R)$

## Transition-based dependency parser

Contains: buffer, stack, tree

Operations:

- Shift transition → Pop buffer, push to stack
- Left arc transition → Dependency from top of stack to second top, remove second top.
- Right arc transition → Dependency from second top of stack to top, remove top.

Terminate when buffer is empty and stack has 1 or less elements

## ***Learn a probabilistic context-free grammar from a treebank***

Estimate rule probabilities → count of specific rule / count of all rules with same left side.

## ***Simulate a transition-based dependency parser***

Stack → ← buffer

[] [I booked a flight from L.A.]

I booked a flight from L.A.

*EVENTUALLY*

[booked] []

```
      _____  
____  |  _  |  ____  ____  
v   |   | v | v |   v |   v
```

I booked a flight from L.A.

## **Semantic Analysis (wildcard)**

### **Word sense**

Lexeme → set of words, same fundamental meaning (run, runs, ran → lexeme RUN)

Lemma → Lexeme you'd put in a lexicon

One lemma → multiple lexemes (word senses)

### **Homonymy**

Same pronunciation/spelling, different meaning

## **Polysemy**

Two senses of a lemma are semantically linked.

## **Synonymy**

When two senses of two different lemmas are (nearly) identical.

If can substitute word A with word B without changing the meaning of the sentence, A & B are synonymous.

## **Antonymy**

Opposite of synonyms. A & B are opposites.

## **Hyponymy**

More specific (car → vehicle, mango → fruit).

Hyponym is the lower word in the word tree.

## **Hypernymy**

Less specific (furniture → chair, fruit → mango).

Hypernym is the upper word in the word tree.

## **WordNet**

Website. Three databases: nouns, verbs, adjectives + adverbs.

Each lemma has synset, a set of one or more senses.

## **Simplified Lesk algorithm**

Given word in a context + number of senses for word.

Textual overlap of non-stopwords between context and sense → score of sense.

## **Word similarity**

How similar is word A to word B? Synonym is boolean relation, want numeric representation.

## **Distributional hypothesis**

Distance between two word senses by finding words with similar distributions in a corpus.

Represent words as vectors.

## Co-occurrence matrix

		context words						
		crown	throne	reign	Sweden	match	goal	play
target words	queen	4	1	1	2	0	0	0
	king	3	2	1	3	1	0	0
	soccer	1	0	0	4	3	4	2
	hockey	0	1	0	1	2	1	1

### ***Simulate the Simplified Lesk algorithm***

Count non-stopword similarities between context and senses, take highest count.

### ***Compute the path length-based similarity of two words***

similarity = (word1, word2) → return  $1 / (1 + \text{pathlength}(\text{word1}, \text{word2}))$ ;

pathlength = number of edges in shortest path between word1 and word2.

pathlength = Basically count the number of words you meet along the way minus the original word.

### ***Derive a co-occurrence matrix from a document collection***

Each cell → number of documents in which target word (row) co-occurs with context word (col).